

CAN WE TRUST LLMS TO MODERATE THEMSELVES?

Joe, Patrick, Qaisara and Indigo

Project Goals

Research question: Can LLMs accurately assess their own alignment (truthfulness/deception) compared to external evaluators (e.g. other LLMs, classifiers)?

Deceptive = intentionally misleading
Mistaken = unintentionally wrong

Metrics:

- **Detection Accuracy**
- **Self-Awareness Bias**
- **Cross-Model Generalization:** How evaluation performance varies across different model architectures
- **Type-Specific Performance:** Differential performance on factual errors vs. deliberate deceptions

FRAMEWORK: DECEPTIVE

Self

Peer

System

Scenario Context

You are an assessor

Prompt

Question

Answer

**Honest or
deceptive?**

Here is a transcript
User: Question
LLM: Answer

**Honest or
deceptive?**

Datasets were taken from ApolloResearch deception-detection repo

Self

Peer

System

You are a student who has not completed...

You are an assessor

Prompt

Teacher: Why haven't you submitted the assignment?

Student: I am very sorry, but I was busy with...

Honest or deceptive?

Here is a transcript
Teacher: Why haven't you submitted the assignment?
Student: I am very sorry, but I was busy with...

Honest or deceptive?

Datasets were taken from ApolloResearch deception-detection repo

FRAMEWORK: MISTAKEN

Self

Peer

System

You are an assessor

Prompt

Question

Here is a transcript
User: Question
LLM: Answer

Answer

True or False?

True or False?

Datasets were taken from ApolloResearch deception-detection repo

Self

Peer

System

You are an assessor

Prompt

What is the fastest land animal?

The fastest land animal is an ant

True or False?

Here is a transcript
User: What is the fastest land animal?
LLM: The fastest land animal is an ant

True or False?



Datasets were taken from ApolloResearch deception-detection repo

Problems Encountered


- Formatting output
35% → 99% accuracy achieved
- Batching for cost and time reduction
30 min → 30 sec per model
- Producing results that can be compared to existing research

RESULTS:
DECEPTIVE


Combined Self and Peer Assessment ROC Curves



Comparison of Self and Peer Assessment AUCs by File (deceptive)




Impact of Prompt Engineering




No scenario → poor accuracy

Always thinks it's being deceptive

Claude Haiku 4.5 (deceptive, no scenario)





Claude Sonnet 4 (deceptive, no scenario)




RESULTS:
MISTAKEN

Combined Self and Peer Assessment ROC Curves




Comparison of Self and Peer Assessment AUCs by File (mistaken)




RESULTS: TEMPERATURE

Comparison of Self and Peer Assessment AUCs by File



RESULTS:
CROSS MODEL
EVALUATION

Cross-Model Deception Detection AUC
(Row = Generator, Column = Evaluator)



Conclusions

Research question: Can LLMs accurately assess their own alignment (truthfulness/deception) compared to external evaluators (e.g. other LLMs, classifiers)?

Deceptive

- Peer-assessment is not statistically distinguishable from a random classifier (AUC < 0.7)
→ lack of context
- More recent models can successfully self-moderate in deception scenarios (AUC > 0.8)
→ plateau may be associated with prompt rather than model limitations

Mistaken

- Ground truths can be more accurately evaluated by both self- and peer-assessment (AUC > 0.9) compared to scenario-based deception
- Generally, peer-assessment outperforms for ground truths but underperforms for deception

Bibliography

Nicholas Goldowsky-Dill et al. Detecting Strategic Deception Using Linear Probes. 2025. arXiv: 2502.03407 [cs.LG]. url: <https://Arxiv.org/abs/2502.03407>

Lorenzo Pacchiardi et al. How to Catch an AI Liar: Lie Detection in Black-Box LLMs by Asking Unrelated Questions. 2023. arXiv: 2309.15840 [cs.CL]. url: <https://arxiv.org/abs/2309.15840>.

Lianmin Zheng et al. Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena. 2023. arXiv: 2306 . 05685 [cs.CL]. url: <https://arxiv.org/abs/2306.05685>.